

Data Structures

Gabriel Istrate

May 19, 2025

How to augment a data structure

- Four steps:
 - 1. Choose underlying data structure.
 - 2. Determine additional information to be maintained.
 - 3. Verify that additional information can be maintained in the D.S. operations.
 - 4. develop new operations required by new fields.

How to augment a data structure (II)

1. Choose red-black trees. Clue: supports other dynamic set operations on total order: MINIMUM, MAXIMUM, SUCCESSOR, PREDECESSOR.
2. We didn't need field size to implement OS-SELECT, OS-RANK, but then operations wouldn't run in $O(\log n)$ time. Additional information to be maintained: sometimes pointer rather than data.
3. Ideally only a few elements need to be updated to maintain D.S. E.g. if we simply stored in each node its rank in the tree then OS-SELECT and OS-UPDATE would be efficient but inserting a smallest node causes changes in the whole tree.
4. Developed OS-SELECT, OS-RANK. Occasionally, instead of new operations, speed-up old ones.

Augmenting red-black trees

Theorem

Let f be a field that augments a RB tree of n nodes, and suppose the contents of f for node x can be computed in $O(1)$ using only information in node x , $\text{left}[x]$ and $\text{right}[x]$, including $f[\text{left}[x]]$ and $f[\text{right}[x]]$. Then we can maintain the values of f in all nodes in T during insertion and deletion without asymptotically affecting $O(\log n)$ performance.

Proof idea: change in field f at a node x propagates only to ancestors of x in the tree.

- closed interval: $[t_1, t_2]$. Also open, half-open intervals.
- $i = [t_1, t_2]$. $\text{low}[i] = t_1$, $\text{high}[i] = t_2$.
- i and i' overlap if $i \cap i' \neq \emptyset$. That is $\text{low}[i] \leq \text{high}[i']$ and $\text{low}[i'] \leq \text{high}[i]$.
- Want: Data structure representing a dynamic set of intervals.
- Must support the following operations:
 - $\text{INTERVAL-INSERT}(T, x)$: adds element x , whose int field contains an interval.
 - $\text{INTERVAL-DELETE}(T, x)$: removes element x from T .
 - $\text{INTERVAL-SEARCH}(T, i)$: return pointer to an element x such that $\text{int}[x]$ overlaps i , or nil if no such element found.

- Any two intervals satisfy **interval trichotomy**: three alternatives:

- i and i' overlap.
- i is to the left of i' ($\text{high}[i] < \text{low}[i']$).
- i is to the right of i' . ($\text{low}[i] > \text{high}[i']$).

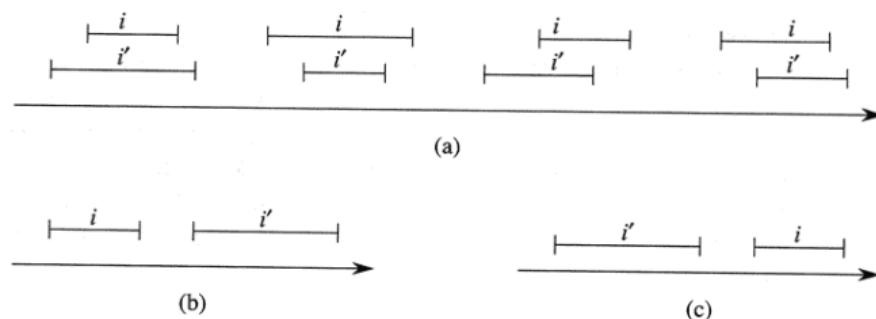


Figure 14.3 The interval trichotomy for two closed intervals i and i' . (a) If i and i' overlap, there are four situations; in each, $\text{low}[i] \leq \text{high}[i']$ and $\text{low}[i'] \leq \text{high}[i]$. (b) The intervals do not overlap, and $\text{high}[i] < \text{low}[i']$. (c) The intervals do not overlap, and $\text{high}[i'] < \text{low}[i]$.

Interval trees: Implementation

1. Possible clue: **intervals (partial) ordering**. Might try to modify a total order. Then **red-black tree**. Each node x stores an interval $\text{int}[x]$.
 - $\text{key}[x] = \text{low}[\text{int}[x]]$.
2. Additional info: $\text{max}[x]$, **the maximum value of any endpoint of an interval stored in the subtree rooted at x** .
3. Maintain info: $\text{max}[x] = \max(\text{high}[\text{int}[x]], \text{max}[\text{left}[x]], \text{max}[\text{right}[x]])$.
4. By applying previous theorem: insertion/deletion $O(\log n)$ while maintaining $\text{max}[x]$.

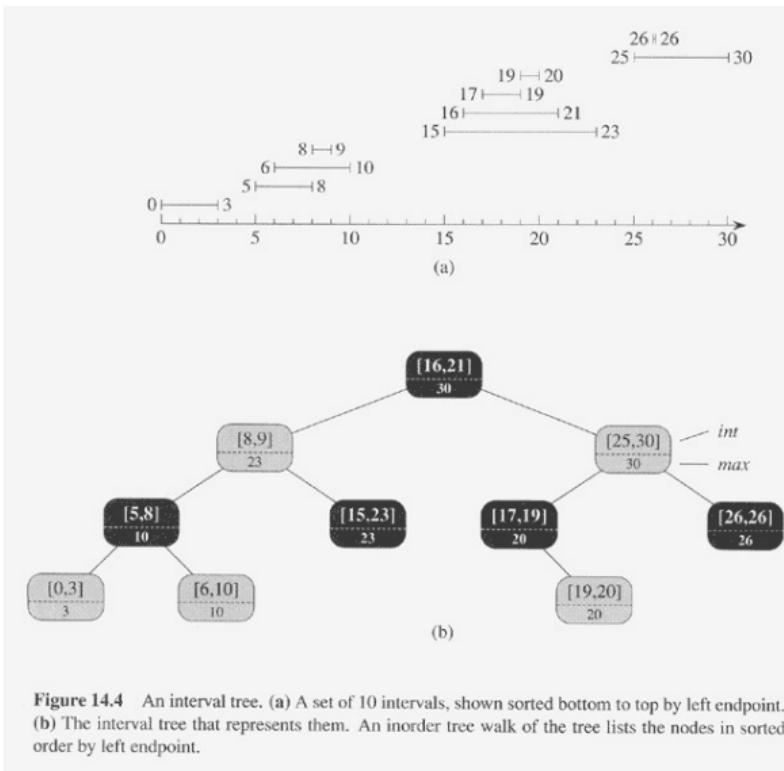


Figure 14.4 An interval tree. (a) A set of 10 intervals, shown sorted bottom to top by left endpoint. (b) The interval tree that represents them. An inorder tree walk of the tree lists the nodes in sorted order by left endpoint.

- finds a node in tree T whose interval overlaps interval i , returns sentinel node $nil[T]$ if no overlapping interval found.
- Search starts at the root and proceeds downwards.
- Chooses $left$ or $right$ subtree based on the maximum element in the left subtree of x .
- If $\max[left[x]] \geq low[i]$ (of course, $left[x] \neq nil[T]$) go left.
- otherwise go right.
- takes $O(\log n)$ time since each basic loop takes $O(1)$ time and the height of the RB tree is $O(\log n)$.

```
INTERVAL-SEARCH( $T, i$ )
1    $x \leftarrow root[T]$ 
2   while  $x \neq nil[T]$  and  $i$  does not overlap  $int[x]$ 
3       do if  $left[x] \neq nil[T]$  and  $max[left[x]] \geq low[i]$ 
4           then  $x \leftarrow left[x]$ 
5           else  $x \leftarrow right[x]$ 
6   return  $x$ 
```

Correctness of INTERVAL-SEARCH

- Why is it enough to examine a single path ?
- Idea: search proceeds in a "safe direction".
- INVARIANT: If tree T contains an interval that overlaps i then there is such an interval in the subtree rooted at x .
- Initialization: clearly satisfied, $x = \text{root}[T]$.
- Either line 4 or line 5 executed.
- Line 5 executed: because $\text{left}[x] = \text{nil}[T]$ or $\text{max}[\text{left}[x]] < \text{low}[i]$. The subtree rooted at $\text{left}[x]$ does not contain any interval that overlaps i .
- If such an interval is found in T , it must be in $\text{right}[x]$.

Correctness of INTERVAL-SEARCH

- Line 4 executed: contrapositive of loop invariant holds.
- If there is no such an interval in the subtree rooted at $\text{left}[x]$ then there is no such interval in tree T .
- Since line 4 executed $\max[\text{left}[x]] \geq \text{low}[i]$. There exists i' with $\text{high}[i'] = \max[\text{left}[x]] \geq \text{low}[i]$.
- i and i' do not overlap, by assumption. By trichotomy $\text{high}[i] < \text{low}[i']$.
- i'' interval in $\text{right}[x]$. Intervals keyed on the low endpoints.
- $\text{high}[i] < \text{low}[i'] \leq \text{low}[i'']$.
- Conclusion: no interval in $\text{right}[x]$ (and thus in T) overlaps i .

Next

- Data Structures for external memory: B-Trees.

Outline:

- Search in secondary storage
- B-Trees
 - ▶ properties
 - ▶ search
 - ▶ insertion

Complexity Model

- Basic assumption so far: *data structures fit completely in main memory (RAM)*
 - ▶ all basic operations have the same cost
 - ▶ even this is a rough approximation, since the main-memory system is not at all “flat”

- Basic assumption so far: *data structures fit completely in main memory (RAM)*
 - ▶ all basic operations have the same cost
 - ▶ even this is a rough approximation, since the main-memory system is not at all “flat”
- However, some applications require more storage than what fits in main memory
 - ▶ we must use data structures that reside in *secondary storage* (i.e., disk)

- Basic assumption so far: *data structures fit completely in main memory (RAM)*
 - ▶ all basic operations have the same cost
 - ▶ even this is a rough approximation, since the main-memory system is not at all “flat”
- However, some applications require more storage than what fits in main memory
 - ▶ we must use data structures that reside in *secondary storage* (i.e., disk)

Disk is 10,000–100,000 times slower than RAM

Idea

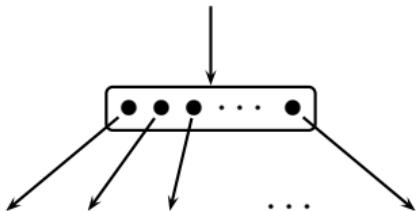
- In a balanced *binary* tree, n keys require a tree of height $h = \lfloor \log_2 n \rfloor$
 - ▶ all the important operations require access to $O(h)$ nodes
 - ▶ each one accounting for *one or very few* basic operations

- In a balanced *binary* tree, n keys require a tree of height $h = \lfloor \log_2 n \rfloor$
 - ▶ all the important operations require access to $O(h)$ nodes
 - ▶ each one accounting for *one or very few* basic operations
- **Idea:** store several keys and pointers to children nodes in a single node

- In a balanced *binary* tree, n keys require a tree of height $h = \lfloor \log_2 n \rfloor$
 - ▶ all the important operations require access to $O(h)$ nodes
 - ▶ each one accounting for *one or very few* basic operations
- **Idea:** store several keys and pointers to children nodes in a single node
 - ▶ in practice we ***increase the degree*** (or *branching factor*) of each node up to $d > 2$, so $h = \lfloor \log_d n \rfloor$
 - ▶ in practice d can be as high as a few thousands

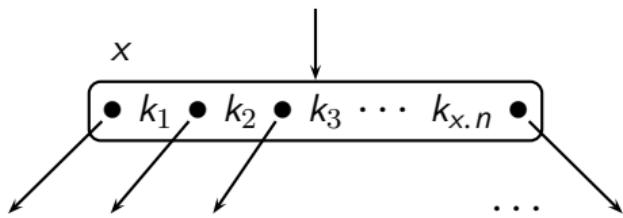
- In a balanced *binary* tree, n keys require a tree of height $h = \lfloor \log_2 n \rfloor$
 - ▶ all the important operations require access to $O(h)$ nodes
 - ▶ each one accounting for *one or very few* basic operations

- **Idea:** store several keys and pointers to children nodes in a single node
 - ▶ in practice we **increase the degree** (or *branching factor*) of each node up to $d > 2$, so $h = \lfloor \log_d n \rfloor$
 - ▶ in practice d can be as high as a few thousands

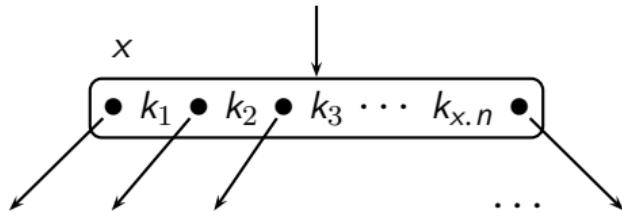


E.g., if $d = 1000$, then
only three accesses ($h = 2$)
cover **up to one billion keys**

Definition of a B-Tree

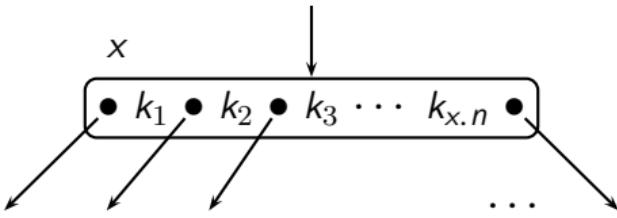


Definition of a B-Tree



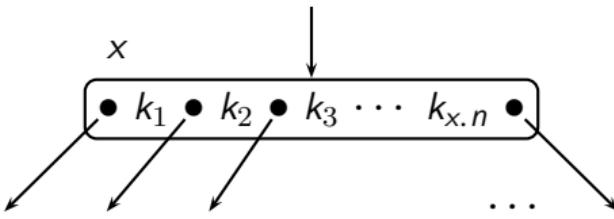
- Every node x has the following fields
 - ▶ $x.n$ is the number of keys stored at each node

Definition of a B-Tree

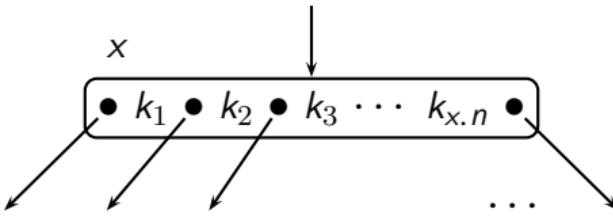


- Every node x has the following fields
 - ▶ $x.n$ is the number of keys stored at each node
 - ▶ $x.key[1] \leq x.key[2] \leq \dots x.key[x.n]$ are the $x.n$ keys stored in nondecreasing order

Definition of a B-Tree

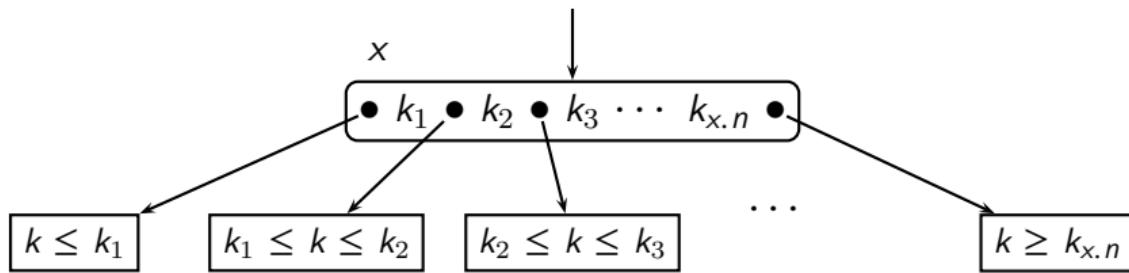


- Every node x has the following fields
 - ▶ $x.n$ is the number of keys stored at each node
 - ▶ $x.key[1] \leq x.key[2] \leq \dots x.key[x.n]$ are the $x.n$ keys stored in nondecreasing order
 - ▶ $x.leaf$ is a Boolean flag that is TRUE if x is a *leaf node* or FALSE if x is an *internal node*

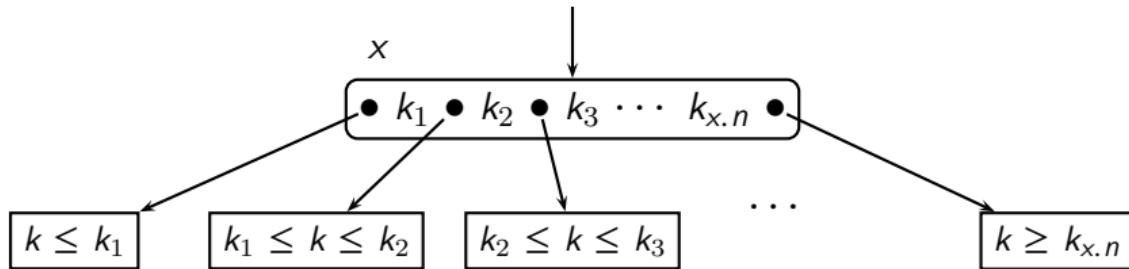


- Every node x has the following fields
 - ▶ $x.n$ is the number of keys stored at each node
 - ▶ $x.key[1] \leq x.key[2] \leq \dots x.key[x.n]$ are the $x.n$ keys stored in nondecreasing order
 - ▶ $x.leaf$ is a Boolean flag that is TRUE if x is a *leaf node* or FALSE if x is an *internal node*
 - ▶ $x.c[1], x.c[2], \dots, x.c[x.n + 1]$ are the $x.n + 1$ pointers to its children, if x is an *internal node*

Definition of a B-Tree (2)

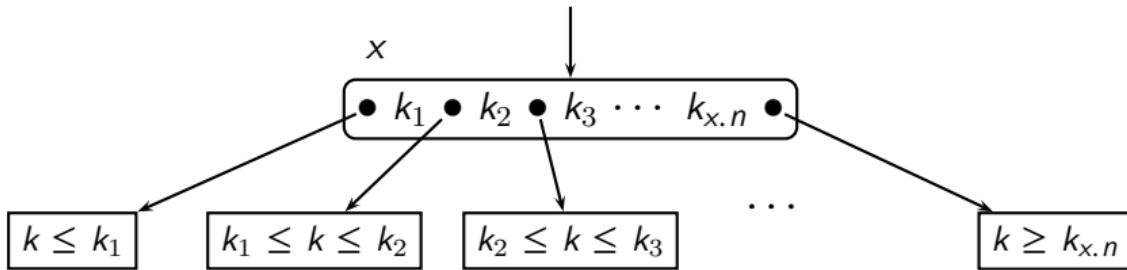


Definition of a B-Tree (2)



- The keys $x.key[i]$ delimit the ranges of keys stored in each subtree

Definition of a B-Tree (2)



- The keys $x.key[i]$ delimit the ranges of keys stored in each subtree

$x.c[1] \rightarrow$ subtree containing keys $k \leq x.key[1]$

$x.c[2] \rightarrow$ subtree containing keys $k, x.key[1] \leq k \leq x.key[2]$

$x.c[3] \rightarrow$ subtree containing keys $k, x.key[2] \leq k \leq x.key[3]$

\dots

$x.c[x.n + 1] \rightarrow$ subtree containing keys $k, k \geq x.key[x.n]$

Definition of a B-Tree (3)

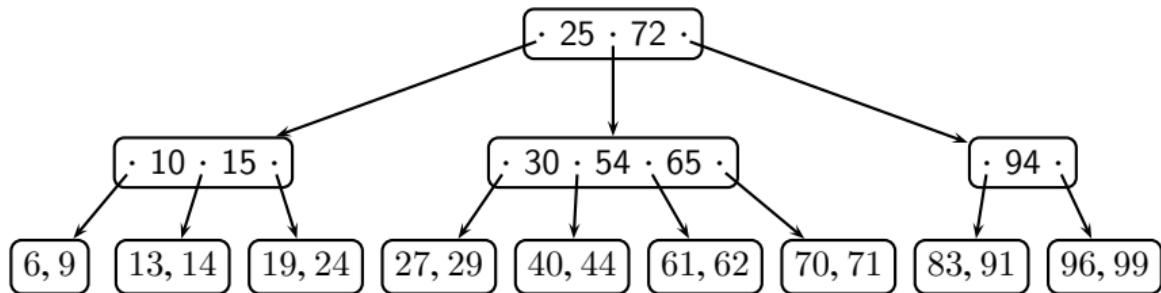
Definition of a B-Tree (3)

- *All leaves have the same depth*

Definition of a B-Tree (3)

- *All leaves have the same depth*
- Let $t \geq 2$ be the **minimum degree** of the B-tree
 - ▶ every node other than the root must have **at least $t - 1$ keys**
 - ▶ every node must contain **at most $2t - 1$ keys**
 - ▶ a node is *full* when it contains exactly $2t - 1$ keys
 - ▶ a full node has $2t$ children

Example



Search in B-Trees

B-Tree-Search(x, k)

```
1   $i = 1$ 
2  while  $i \leq x.n$  and  $k > x.key[i]$ 
3       $i = i + 1$ 
4  if  $i \leq x.n$  and  $k == x.key[i]$ 
5      return  $(x, i)$ 
6  if  $x.leaf$ 
7      return NIL
8  else Disk-Read( $x.c[i]$ )
9      return B-Tree-Search( $x.c[i], k$ )
```

Height of a B-Tree

Height of a B-Tree

- **Theorem:** the height of a B-tree containing $n \geq 1$ keys and with a minimum degree $t \geq 2$ is

$$h \leq \log_t \frac{n+1}{2}$$

Height of a B-Tree

- **Theorem:** the height of a B-tree containing $n \geq 1$ keys and with a minimum degree $t \geq 2$ is

$$h \leq \log_t \frac{n+1}{2}$$

Proof:

- ▶ $n \geq 1$, so the root has at least one key (and therefore two children)

Height of a B-Tree

- **Theorem:** the height of a B-tree containing $n \geq 1$ keys and with a minimum degree $t \geq 2$ is

$$h \leq \log_t \frac{n+1}{2}$$

Proof:

- ▶ $n \geq 1$, so the root has at least one key (and therefore two children)
- ▶ every other node has at least t children

Height of a B-Tree

- **Theorem:** the height of a B-tree containing $n \geq 1$ keys and with a minimum degree $t \geq 2$ is

$$h \leq \log_t \frac{n+1}{2}$$

Proof:

- ▶ $n \geq 1$, so the root has at least one key (and therefore two children)
- ▶ every other node has at least t children
- ▶ in the worst case, there are two subtrees (of the root) each one containing a total of $(n - 1)/2$ keys, and each one consisting of t -degree nodes, with each node containing $t - 1$ keys

- **Theorem:** the height of a B-tree containing $n \geq 1$ keys and with a minimum degree $t \geq 2$ is

$$h \leq \log_t \frac{n+1}{2}$$

Proof:

- ▶ $n \geq 1$, so the root has at least one key (and therefore two children)
- ▶ every other node has at least t children
- ▶ in the worst case, there are two subtrees (of the root) each one containing a total of $(n - 1)/2$ keys, and each one consisting of t -degree nodes, with each node containing $t - 1$ keys
- ▶ each subtree contains $1 + t + t^2 + \dots + t^{h-1}$ nodes, each one containing $t - 1$ keys

- **Theorem:** the height of a B-tree containing $n \geq 1$ keys and with a minimum degree $t \geq 2$ is

$$h \leq \log_t \frac{n+1}{2}$$

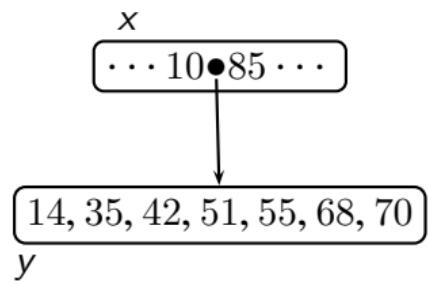
Proof:

- ▶ $n \geq 1$, so the root has at least one key (and therefore two children)
- ▶ every other node has at least t children
- ▶ in the worst case, there are two subtrees (of the root) each one containing a total of $(n - 1)/2$ keys, and each one consisting of t -degree nodes, with each node containing $t - 1$ keys
- ▶ each subtree contains $1 + t + t^2 + \dots + t^{h-1}$ nodes, each one containing $t - 1$ keys, so

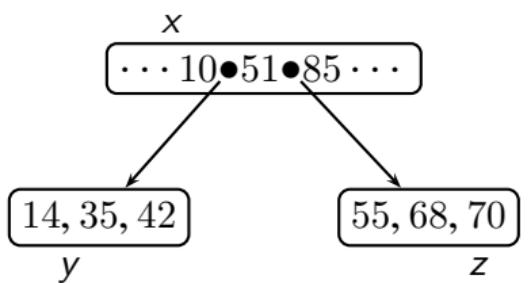
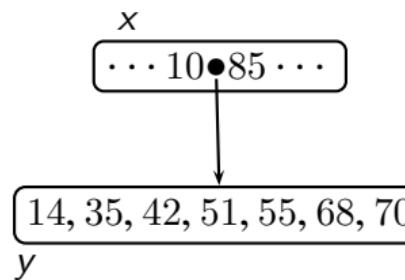
$$n \geq 1 + 2(t^h - 1)$$

Splitting

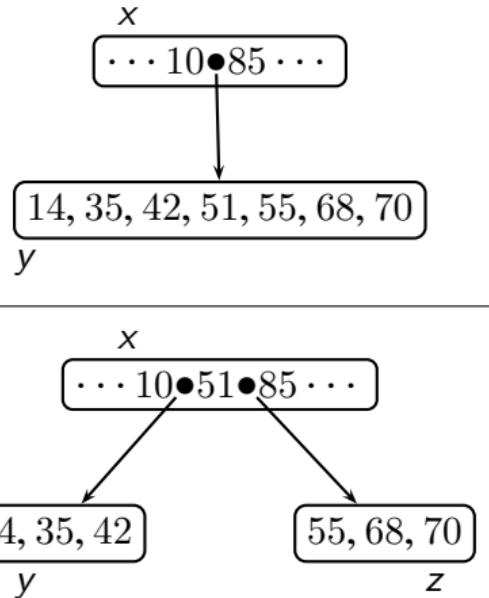
Splitting



Splitting



Splitting



B-Tree-Split-Child(x, i, y)

```
1   $z = \text{Allocate-Node}()$ 
2   $z.\text{leaf} = y.\text{leaf}$ 
3   $z.n = t - 1$ 
4  for  $j = 1$  to  $t - 1$ 
5       $z.\text{key}[j] = y.\text{key}[j + t]$ 
6  if not  $y.\text{leaf}$ 
7      for  $j = 1$  to  $t$ 
8           $z.c[j] = y.c[j + t]$ 
9   $y.n = t - 1$ 
10 for  $j = x.n + 1$  downto  $i + 1$ 
11      $x.c[j + 1] = x.c[j]$ 
12 for  $j = x.n$  downto  $i$ 
13      $x.key[j + 1] = x.key[j]$ 
14  $x.key[i] = y.key[t]$ 
15  $x.n = x.n + 1$ 
16 Disk-Write( $y$ )
17 Disk-Write( $z$ )
18 Disk-Write( $x$ )
```

Complexity of B-Tree-Split-Child

- What is the complexity of **B-Tree-Split-Child**?

Complexity of B-Tree-Split-Child

- What is the complexity of **B-Tree-Split-Child**?
- $\Theta(t)$ basic CPU operations

Complexity of B-Tree-Split-Child

- What is the complexity of **B-Tree-Split-Child**?
- $\Theta(t)$ basic CPU operations
- 3 **Disk-Write** operations

B-Tree-Split-Child(x, i, y)

```
1  z = Allocate-Node()
2  z.leaf = y.leaf
3  z.n = t - 1
4  for j = 1 to t - 1
5      x.key[j] = x.key[j + t]
6  if not x.leaf
7      for j = 1 to t
8          z.c[j] = y.c[j + t]
9  y.n = t - 1
10 for j = x.n + 1 downto i + 1
11     x.c[j + 1] = x.c[j]
12 for j = x.n downto i
13     x.key[j + 1] = x.key[j]
14 x.key[i] = y.key[t]
15 x.n = x.n + 1
16 Disk-Write(y)
17 Disk-Write(z)
18 Disk-Write(x)
```

Insertion Under Non-Full Node

Insertion Under Non-Full Node

B-Tree-Insert-Nonfull(x, k)

```
1   $i = x.n$                                 // assume  $x$  is not full
2  if  $x.\text{leaf}$ 
3      while  $i \geq 1$  and  $k < x.\text{key}[i]$ 
4           $x.\text{key}[i + 1] = x.\text{key}[i]$ 
5           $i = i - 1$ 
6           $x.\text{key}[i + 1] = k$ 
7           $x.n = x.n + 1$ 
8          Disk-Write( $x$ )
9  else while  $i \geq 1$  and  $k < x.\text{key}[i]$ 
10          $i = i - 1$ 
11          $i = i + 1$ 
12         Disk-Read( $x.c[i]$ )
13         if  $x.c[i].n == 2t - 1$       // child  $x.c[i]$  is full
14             B-Tree-Split-Child( $x, i, x.c[i]$ )
15             if  $k > x.\text{key}[i]$ 
16                  $i = i + 1$ 
17             B-Tree-Insert-Nonfull( $x.c[i], k$ )
```

Insertion Procedure

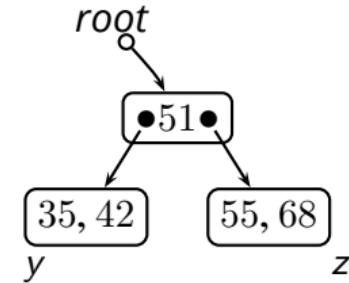
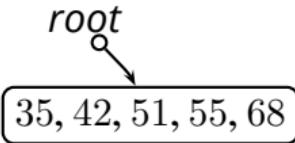
B-Tree-Insert(T, k)

```
1   $r = T.root$ 
2  if  $r.n == 2t - 1$ 
3       $s = \text{Allocate-Node}()$ 
4       $T.root = s$ 
5       $s.leaf = \text{FALSE}$ 
6       $s.n = 0$ 
7       $s.c[1] = r$ 
8      B-Tree-Split-Child( $s, 1, r$ )
9      B-Tree-Insert-Nonfull( $s, k$ )
10 else B-Tree-Insert-Nonfull( $r, k$ )
```

Insertion Procedure

B-Tree-Insert(T, k)

```
1   $r = T.root$ 
2  if  $r.n == 2t - 1$ 
3       $s = \text{Allocate-Node}()$ 
4       $T.root = s$ 
5       $s.leaf = \text{FALSE}$ 
6       $s.n = 0$ 
7       $s.c[1] = r$ 
8      B-Tree-Split-Child( $s, 1, r$ )
9      B-Tree-Insert-Nonfull( $s, k$ )
10 else B-Tree-Insert-Nonfull( $r, k$ )
```



Complexity of Insertion

- What is the complexity of **B-Tree-Insert**?

Complexity of Insertion

- What is the complexity of **B-Tree-Insert**?
- $O(th) = O(t \log_t n)$ basic CPU steps operations

Complexity of Insertion

- What is the complexity of **B-Tree-Insert**?
- $O(th) = O(t \log_t n)$ basic CPU steps operations
- $O(h) = O(\log_t n)$ disk-access operations

Complexity of Insertion

- What is the complexity of **B-Tree-Insert**?
- $O(th) = O(t \log_t n)$ basic CPU steps operations
- $O(h) = O(\log_t n)$ disk-access operations
- The best value for t can be determined according to
 - ▶ the ratio between CPU (RAM) speed and disk-access time
 - ▶ the *block-size* of the disk, which determines the maximum size of an object that can be accessed (read/write) in one shot

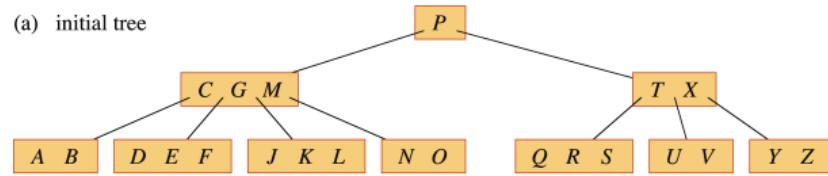
- Not hard, but more tedious.

- Not hard, but more tedious.
- Main problem: have to delete keys from every nodes, not just leaves

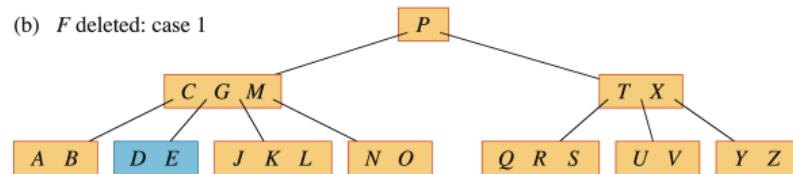
- Not hard, but more tedious.
- Main problem: have to delete keys from every nodes, not just leaves
- Insertion: merging nodes. Deletion: splitting nodes

- Not hard, but more tedious.
- Main problem: have to delete keys from every nodes, not just leaves
- Insertion: merging nodes. Deletion: splitting nodes
- $O(h)$ disk operations. $O(th) = O(t \cdot \log_t(n))$ CPU time.

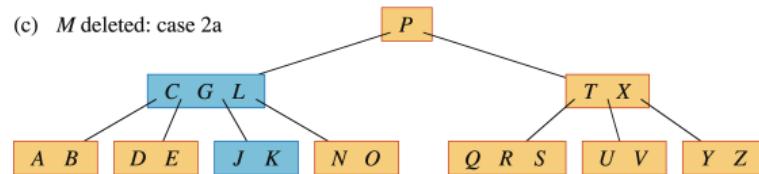
(a) initial tree



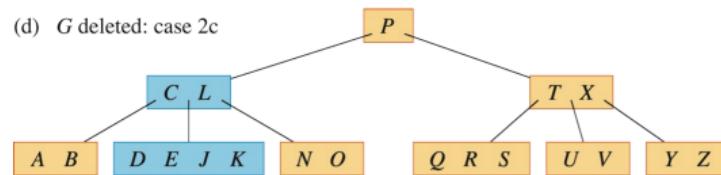
(b) F deleted: case 1



(c) M deleted: case 2a



(d) G deleted: case 2c



Want to know more?

Foundations and Trends® In
Theoretical Computer Science
2:4

Algorithms and Data Structures for External Memory

Jeffrey Scott Vitter